

UNIVERSITY OF SCIENCE, VNU-HCM

FACULTY OF INFORMATION TECHNOLOGY
ADVANCED PROGRAM IN COMPUTER SCIENCE



CS486 - Introduction to Database System

Phase 2 Report

Campus Space Management System

Students:

Tran Canh Anh Tuan 24125048
Le Dat Phuc An 24125052
Vong Chi Van 24125049
Pham Nguyen Minh Quan
24125041

Lecturers:

Nguyen Ngoc Toan, MSc.
Nguyen Ngoc Minh Chau, MSc.
Le Thi Nhan, PhD

Summary

1	Introduction	3
1.1	Background and Phase 2 Objectives	3
1.2	Group Members and Contributions	3
1.3	Phase 2 Deliverables	4
2	Requirement Changes and Database Evolution	4
2.1	Maintenance Impact Levels	4
2.2	Affected Concepts and Rules	5
2.3	Updated Conceptual Overview	5
2.4	Schema Migration	7
2.5	Integrity Boundary	7
3	Concurrency Design and Test Results	8
3.1	Identified Conflict	8
3.2	Protected Approval Protocol	8
3.3	Unsafe Test	9
3.4	Protected Test	10
4	Generated Data and Required Analytical Queries	11
4.1	Data Generation	11
4.2	Query 01: Approved Booking Hours per Space	12
4.3	Query 02: Approved Bookings by Weekday and Hour	12
4.4	Query 03: Available-Space Finder	13
4.5	Query 04: Bookings Affected by Maintenance Escalation	14
4.6	Requirement Traceability	14
5	Indexing and Query Tuning	15
5.1	Scope and Methodology	15
5.2	Selected Indexes	15
5.3	Before-and-After Results	15
5.4	Execution-Plan Interpretation	16
5.5	Trade-offs and Recommendation	16
6	Normalization Validation	17
6.1	Functional Dependencies and Candidate Keys	17
6.2	Normal-Form Proof	17
7	LLM and Agent Improvement Process	18
7.1	Model and Framework	18
7.2	Improvement Cycle	18
8	Conclusion	19

§1 Introduction

§1.1 Background and Phase 2 Objectives

Phase 2 extends the Campus Space Management System after a one-semester pilot. The original database treated every active maintenance record as a complete booking block. The revised requirements distinguish disruptive maintenance from advisory work, introduce automatic approval for selected role–space-type combinations, require correctness under concurrent approval, and add analytical reporting over a substantially larger history. The implementation therefore has five objectives:

- represent advisory and out-of-service maintenance separately;
- preserve proof that requesters were notified about active advisories;
- prevent overlapping approvals even when automatic and staff workflows run concurrently;
- support four required operational and analytical reports; and
- demonstrate measurable indexing improvements on at least 100,000 bookings.

§1.2 Group Members and Contributions

Member	Student ID	Phase 2 contributions
Tran Canh Anh Tuan	24125048	Shared updated ERD and logical design (09); requirement-change analysis and concurrency design, implementation, and tests (08, 11–13).
Le Dat Phuc An	24125052	Shared updated ERD and logical design (09); schema migration, data generation, and analytical queries (10, 14, 16).
Vong Chi Van	24125049	Shared updated ERD and logical design (09); requirement-change analysis and concurrency design, implementation, and tests (08, 11–13).
Pham Nguyen Minh Quan	24125041	Shared updated ERD and logical design (09); index tuning and analytical-query refinement (15, 16).

Table 1: Group membership and deliverable ownership.

§1.3 Phase 2 Deliverables

Output	Purpose
08	Requirement-change and concurrency-conflict analysis
09	Updated ERD and logical schema
10	Data-preserving migration from the Phase 1 schema
11--13	Concurrency design, protected procedure, and two-session tests
14	Deterministic three-academic-year data generator
15	Indexing experiment and before/after measurements
16	Four required analytical queries considered in this report

Table 2: Phase 2 artifact summary.

This report is based on the submitted artifacts and the captured SQL Server test evidence. SQL excerpts are intentionally shortened; the complete executable statements remain in the `outputs/` directory.

§2 Requirement Changes and Database Evolution

§2.1 Maintenance Impact Levels

The revised rule separates maintenance into two impact levels:

`out_of_service`

The space is unusable. An active record blocks every booking whose half-open interval overlaps the maintenance interval.

`advisory`

The space remains bookable. Every active advisory at booking time must be communicated to the requester, and the database must retain evidence of that communication.

Several active records can coexist for one space. An open record can also be escalated or downgraded. Escalation to `out_of_service` does not erase earlier approvals; instead, staff must be able to retrieve the affected bookings and contact their requesters.

§2.2 Affected Concepts and Rules

Concept	Phase 2 change
MAINTENANCE_RECORD	Adds <code>impact_level</code> ; overlapping active out-of-service work blocks approval, while advisory work permits approval and requires notification.
BOOKING_NOTIFICATION	Associates a booking and maintenance record with a notification type and timestamp, preserving advisory and escalation history.
BOOKING_DECISION	Replaces the approval-only relation and represents automatic and staff decisions uniformly.
ROLE, SPACE_TYPE	Normalize policy dimensions that were stored as text in Phase 1.
AUTO_USAGE_POLICY	Defines which role and space-type combinations are eligible for automatic processing.
Approval workflow	Every approval path must use the same atomic availability-check-and-write protocol.

Table 3: Principal requirement-driven design changes.

§2.3 Updated Conceptual Overview

The Phase 2 design contains eleven relations. Figure 1 emphasizes foreign-key dependencies and the common data flow from requests to decisions, sessions, maintenance notifications, and reports.




Figure 1: Updated Schema Overview

Relation	Primary key	Important foreign/candidate keys
ROLE	role_id	role_name is unique
USERS	user_id	role_id; email is unique
SPACE_TYPE	space_type_id	space_type_name is unique
SPACES	space_code	space_type_id; (building, room_number) is unique
AUTO_USAGE_POLICY	(space_type_id, role_id)	both columns are foreign keys
FACILITY	facility_id	space_code
BOOKING_REQUEST	booking_id	user_id, space_code
BOOKING_DECISION	decision_id	booking_id is unique; decided_by_staff
USAGE_SESSION	session_id	decision_id is unique; two staff references
MAINTENANCE_RECORD	maintenance_id	space_code and three user-role references
BOOKING_NOTIFICATION	(booking_id, maintenance_id, notification_type)	booking_id and maintenance_id

Table 4: Updated relation and key summary.

§2.4 Schema Migration

Output 10 applies the changes on top of the Phase 1 database rather than rebuilding it. Textual roles and space types are first copied into lookup tables; the original rows then receive foreign-key identifiers. Approval records are transformed into booking decisions before the old approval table is removed. Usage sessions are reconnected through their decisions, and new policy and notification relations begin empty because Phase 1 contained no equivalent facts.

Migration operation	Preservation approach
Normalize roles and space types	Insert distinct Phase 1 values, populate new identifiers by joins, enforce non-null foreign keys, then remove text columns.
Replace booking approval	Insert one <code>BOOKING_DECISION</code> for each existing approval, preserving its identifier, booking, staff member, reason, and time.
Reconnect usage sessions	Resolve each session's booking to the migrated decision, enforce a unique decision reference, and remove the former booking reference.
Extend maintenance	Add <code>impact_level</code> with <code>out_of_service</code> as the migration default, preserving the stricter Phase 1 behavior.
Add new concepts	Create empty automatic-policy and booking-notification relations because no reliable Phase 1 facts exist to infer.
Verify	Finish with table row counts; expected migrated Phase 1 counts include 20 users, 12 spaces, 30 requests, 25 decisions, 16 sessions, and 12 maintenance records.

Table 5: Data-preserving migration strategy.

§2.5 Integrity Boundary

Primary keys, foreign keys, uniqueness, and check constraints enforce structural rules such as valid identifiers and time order. Availability, interval overlap, maintenance impact, automatic-policy eligibility, and advisory notification require multiple rows or relations. Those rules are therefore enforced inside transactional procedures rather than by declarative constraints alone.

§3 Concurrency Design and Test Results

§3.1 Identified Conflict

The critical anomaly is a check-then-act race. Consider two pending requests for the same space:

$$A = [09:00, 11:00), \quad B = [10:00, 12:00).$$

Without a shared serialization resource, Session A and Session B can both read *no approved overlap* before either transaction inserts its decision. The unique constraint on `BOOKING_DECISION.booking_id` cannot prevent this outcome because the sessions decide two different booking identifiers.

Step	Session A	Session B
1	Checks overlap; none exists	
2		Checks overlap; none exists
3	Inserts approval A	Inserts approval B
4	Commits	Commits

Table 6: Unsafe concurrent check-and-write schedule.

The required invariant uses half-open intervals: two approved, non-cancelled bookings for one `space_code` must never satisfy

$$existing.start < candidate.end \quad \wedge \quad candidate.start < existing.end.$$

§3.2 Protected Approval Protocol

The protected approval procedure uses pessimistic SQL Server locking in one transaction:

1. Read the target `BOOKING_REQUEST` with `UPDLOCK`. This keeps its space, interval, and pending status stable.
2. Lock the existing `SPACES` row for that `space_code` with `UPDLOCK`.
3. Check approved, non-cancelled overlapping bookings.
4. Insert the decision and update the request to `approved`.
5. Commit, releasing both update locks.

Listing 1: Core locking order in the protected approval procedure.

```

1 BEGIN TRANSACTION;
2
3 SELECT @space_code = br.space_code,
4         @start_time = br.start_time,
5         @end_time = br.end_time,
6         @status = br.status
7 FROM dbo.BOOKING_REQUEST AS br WITH (UPDLOCK)
8 WHERE br.booking_id = @booking_id;
```



```

9
10 SELECT @space_lock_code = s.space_code
11 FROM dbo.SPACES AS s WITH (UPDLOCK)
12 WHERE s.space_code = @space_code;
13
14 -- overlap check, decision insert, and status update
15 COMMIT TRANSACTION;

```

All approval routes, including automatic approval, call the same procedure. Competing requests for the same space therefore request incompatible update locks on one common row. Requests for different spaces lock different rows and can still proceed concurrently. The fixed order `BOOKING_REQUEST` then `SPACES` also reduces deadlock risk.

§3.3 Unsafe Test

The unsafe scripts deliberately bypass the protected procedure. Session A pauses after its availability check, Session B performs the same check and commits, and Session A then commits its stale decision. Figure 2 records the resulting violation.

The screenshot shows a SQL Server Enterprise Manager interface. The top pane displays a script named '04-verify-unsafe.sql' with the following content:

```

-- Expected: both unsafe bookings are approved and form one overlapping pair.
USE campus_space_management;
GO
SET NOCOUNT ON;
SELECT
    br.booking_id,
    br.space_code,
    br.start_time,
    br.end_time,
    br.status,
    d.is_approved
FROM dbo.BOOKING_REQUEST AS br
LEFT JOIN dbo.BOOKING_DECISION AS d
ON d.booking_id = br.booking_id

```

The bottom pane shows the results of the query. It contains two tables. The first table, 'Results', has columns: booking_id, space_code, start_time, end_time, status, and is_approved. It contains two rows:

booking_id	space_code	start_time	end_time	status	is_approved
G08_UNSAFE_B1	G08_CONC_ROOM	2035-03-10 09:00:00.000	2035-03-10 11:00:00.000	approved	1
G08_UNSAFE_B2	G08_CONC_ROOM	2035-03-10 10:00:00.000	2035-03-10 12:00:00.000	approved	1

The second table, 'Messages', has columns: approved_booking_count and overlapping_approved_pair_count. It contains one row:

approved_booking_count	overlapping_approved_pair_count
2	1


Figure 2: Unsafe verification: both requests are approved and one overlapping approved pair exists.

The verification reports

$$approved_booking_count = 2, \quad overlapping_approved_pair_count = 1.$$

§3.4 Protected Test

In the protected run, Session A acquires both update locks and holds them for twelve seconds. Figure 3 shows its successful decision.




```
1  -- 06-safe-session-A.sql
2  -- Run first in Window A.
3  -- During the 12-second hold, run 07-safe-session-B.sql in Window B.
4
5  USE campus_space_management;
6  GO
7
8  EXEC dbo.usp_G08_ApproveBookingConcurrentSafe
9      @booking_id = 'G08_SAFE_B1',
10     @is_automatic = 0,
11     @decided_by_staff = 'G08_STAFF_1',
12     @decision_reason = 'Protected Session A',
13     @test_hold_seconds = 12;
14 GO
15
```

88 % No issues found Ln: 11, Ch: 39 SPC CRLF UTF-8

Results	Messages								
1	<table border="1"><thead><tr><th>decision_id</th><th>booking_id</th><th>space_code</th><th>result</th></tr></thead><tbody><tr><td>G08D5E69359A560A4212</td><td>G08_SAFE_B1</td><td>G08_CONC_ROOM</td><td>approved</td></tr></tbody></table>	decision_id	booking_id	space_code	result	G08D5E69359A560A4212	G08_SAFE_B1	G08_CONC_ROOM	approved
decision_id	booking_id	space_code	result						
G08D5E69359A560A4212	G08_SAFE_B1	G08_CONC_ROOM	approved						

Figure 3: Protected Session A approves the first request while holding the booking and space locks.

Session B can lock its distinct request row but must wait for the same SPACES row. After Session A commits, Session B rechecks the database, detects the committed overlap, and receives application error 52105, as shown in Figure 4.



```
1  -- 07-safe-session-B.sql
2  -- Run in Window B while Session A holds the BOOKING_REQUEST and SPACES
3  -- update locks.
4
5  USE campus_space_management;
6  GO
7
8  BEGIN TRY
9      EXEC dbo.usp_G08_ApproveBookingConcurrentSafe
10         @booking_id = 'G08_SAFE_B2',
11         @is_automatic = 0,
12         @decided_by_staff = 'G08_STAFF_1',
13         @decision_reason = 'Protected Session B',
14         @test_hold_seconds = 0;
15
16      THROW 52610,
17          'UNEXPECTED: Session B was approved instead of detecting the overlap.',
18          1;
19
```

88 % No issues found Ln: 22, Ch: 31 SPC CRLF UT

Results	Messages						
1	<table border="1"><thead><tr><th>error_number</th><th>error_message</th><th>result</th></tr></thead><tbody><tr><td>52105</td><td>Another approved booking overlaps this space and time.</td><td>EXPECTED: Session B waited for the SPACES update lock and then detected the overlap.</td></tr></tbody></table>	error_number	error_message	result	52105	Another approved booking overlaps this space and time.	EXPECTED: Session B waited for the SPACES update lock and then detected the overlap.
error_number	error_message	result					
52105	Another approved booking overlaps this space and time.	EXPECTED: Session B waited for the SPACES update lock and then detected the overlap.					

Figure 4: Protected Session B receives the expected overlap error after waiting for the space lock.

The final assertion in Figure 5 confirms that the first request is approved, the second remains pending, and no overlapping approved pair exists.

```

1  -- 08-verify-safe.sql
2  -- Expected: exactly one safe booking is approved and no approved overlap exists.
3
4  USE campus_space_management;
5  GO
6
7  SET NOCOUNT ON;
8
9  SELECT
10     br.booking_id,
11     br.space_code,
12     br.start_time,
13     br.end_time,
14     br.status,
15     d.is_approved
16 FROM   dbo.BOOKING_REQUEST AS br
17 LEFT JOIN  dbo.BOOKING_DECISION AS d
18   ON d.booking_id = br.booking_id
19 WHERE br.booking_id IN (
20     'G08_SAFE_B1',
21     'G08_SAFE_B2'
22 )
23 ORDER BY br.booking_id;
24
25 DECLARE @approved_count INT = (
26     SELECT COUNT(*)

```

88 % No issues found Ln: 15, Ch: 18 SPC CRLF UT

booking_id	space_code	start_time	end_time	status	is_approved
G08_SAFE_B1	G08_CONC_ROOM	2035-03-11 09:00:00.000	2035-03-11 11:00:00.000	approved	1
G08_SAFE_B2	G08_CONC_ROOM	2035-03-11 10:00:00.000	2035-03-11 12:00:00.000	pending	NULL

approved_booking_count	overlapping_approved_pair_count
1	0

Figure 5: Final protected-state verification.

Test	Approved bookings	Overlapping approved pairs
Unsafe workflow	2	1
Protected workflow	1	0

Table 7: Observed concurrency-test results.

The test demonstrates both sides of the claim: the unprotected check can violate the invariant, while the common per-space update lock serializes the check-and-write section and prevents the violation.

§4 Generated Data and Required Analytical Queries

§4.1 Data Generation

Output 14 is a deterministic and rerunnable Microsoft SQL Server generator. It creates exactly 100,000 booking requests in a reserved identifier namespace while preserving migrated Phase 1 rows. The generated period is 1 September 2023 through 31 August 2026, covering three complete academic years.

Coverage dimension	Generated evidence
Scale	100,000 generated requests; the tuning database contains 100,030 requests after retaining Phase 1 data.
Users and spaces	2,000 requesters, three processing/staff accounts, and 50 generated spaces.
Booking lifecycle	Approved, rejected, pending, completed, cancelled, and no-show cases.
Processing modes	Automatic and staff decisions consistent with <code>AUTO_USAGE_POLICY</code> .
Maintenance	Advisory and out-of-service records, including overlapping active work and notification records.
Integrity checks	Assertions verify the row count, date span, lifecycle coverage, policy consistency, absence of approved overlaps, and advisory-notification completeness.

Table 8: Generated dataset characteristics.

The performance experiment reports 100,030 booking requests, 98,025 booking decisions, 782 maintenance records, 182 facilities, 62 spaces, and six space types. This scale is sufficient to make scan-versus-seek differences visible.

§4.2 Query 01: Approved Booking Hours per Space

Business question. What is the total approved booking time for every space during a supplied semester?

Users. Facility Manager.

The query joins requests to approved decisions, spaces, and space types. It excludes later cancellations, applies a half-open semester boundary, and aggregates duration in minutes before converting it to decimal hours.

Listing 2: Core of Query 01.

```

1 SELECT s.space_code, s.space_name, st.space_type_name,
2        COUNT(*) AS approved_booking_count,
3        CAST(SUM(DATEDIFF(MINUTE, br.start_time, br.end_time))
4             / 60.0 AS DECIMAL(10,2)) AS total_approved_hours
5 FROM BOOKING_REQUEST AS br
6 JOIN BOOKING_DECISION AS bd ON bd.booking_id = br.booking_id
7 JOIN SPACES AS s ON s.space_code = br.space_code
8 JOIN SPACE_TYPE AS st ON st.space_type_id = s.space_type_id
9 WHERE bd.is_approved = 1
10    AND br.status <> 'cancelled'
11    AND br.start_time >= @SemesterStart
12    AND br.start_time < @SemesterEnd
13 GROUP BY s.space_code, s.space_name, st.space_type_name;
```

§4.3 Query 02: Approved Bookings by Weekday and Hour

Business question. How are approved booking decisions distributed by weekday and hour during a supplied semester?

Users. Facility Manager and Department Administrator.

The implementation uses `BOOKING_DECISION.decision_time`, filters an optional hour interval, and groups by the SQL Server weekday name, weekday number, and hour. Cancelled bookings are excluded even if their historical decision was approved.

Listing 3: Core of Query 02.

```
1 SELECT DATENAME(WEEKDAY, bd.decision_time) AS weekday_name,
2        DATEPART(WEEKDAY, bd.decision_time) AS weekday_number,
3        DATEPART(HOUR, bd.decision_time) AS decision_hour,
4        COUNT(*) AS booking_count
5 FROM BOOKING_DECISION AS bd
6 JOIN BOOKING_REQUEST AS br ON br.booking_id = bd.booking_id
7 WHERE bd.is_approved = 1
8        AND br.status <> 'cancelled'
9        AND bd.decision_time >= @SemesterStart
10       AND bd.decision_time < @SemesterEnd
11       AND DATEPART(HOUR, bd.decision_time)
12          BETWEEN @FromHour AND @ToHour
13 GROUP BY DATENAME(WEEKDAY, bd.decision_time),
14          DATEPART(WEEKDAY, bd.decision_time),
15          DATEPART(HOUR, bd.decision_time);
```

§4.4 Query 03: Available-Space Finder

Business question. Which available spaces meet a required capacity and complete facility list during a requested time interval?

Users. Students, lecturers, teaching assistants, and facility staff.

A result must meet all five conditions: operational status, adequate capacity, every requested facility, no approved non-cancelled overlap, and no active out-of-service maintenance overlap. Advisory maintenance does not remove a result; it is returned as a warning flag.

Listing 4: Key predicates in Query 03.

```
1 WHERE s.current_status = 'available'
2        AND s.capacity >= @MinCapacity
3        AND @RequiredFacilityCount = (
4            SELECT COUNT(DISTINCT f.facility_name)
5            FROM FACILITY AS f
6            JOIN @RequiredFacilities AS rf
7              ON rf.facility_name = f.facility_name
8            WHERE f.space_code = s.space_code
9        )
10       AND NOT EXISTS (
11           SELECT 1
12           FROM BOOKING_REQUEST AS br
13           JOIN BOOKING_DECISION AS bd
14             ON bd.booking_id = br.booking_id
15           WHERE br.space_code = s.space_code
16                 AND bd.is_approved = 1
17                 AND br.status <> 'cancelled'
18                 AND br.start_time < @RequestedEnd
19                 AND br.end_time > @RequestedStart
20       )
21       AND NOT EXISTS (
22           SELECT 1
23           FROM MAINTENANCE_RECORD AS m
```

```

24     WHERE m.space_code = s.space_code
25           AND m.impact_level = 'out_of_service'
26           AND m.status IN ('pending', 'in_progress')
27           AND m.start_time < @RequestedEnd
28           AND (m.end_time IS NULL
29                OR m.end_time > @RequestedStart)
30 );

```

§4.5 Query 04: Bookings Affected by Maintenance Escalation

Business question. Which previously approved bookings were affected when a selected advisory record was escalated to out-of-service?

Users. Facility Staff and Facility Manager.

The notification table supplies durable evidence of an actual escalation. The report selects OUT_OF_SERVICE notifications for the chosen maintenance identifier, requires an approved decision, and returns requester contact details, booking times, decision time, escalation time, and maintenance details.

Listing 5: Core of Query 04.

```

1 SELECT br.booking_id, u.full_name, u.email,
2        br.space_code, br.start_time, br.end_time,
3        br.status AS booking_status,
4        bd.decision_time,
5        bn.notification_time AS escalation_time,
6        m.problem_description, m.impact_level
7 FROM BOOKING_NOTIFICATION AS bn
8 JOIN BOOKING_REQUEST AS br ON br.booking_id = bn.booking_id
9 JOIN BOOKING_DECISION AS bd ON bd.booking_id = br.booking_id
10 JOIN MAINTENANCE_RECORD AS m
11     ON m.maintenance_id = bn.maintenance_id
12 JOIN USERS AS u ON u.user_id = br.user_id
13 WHERE bn.maintenance_id = @MaintenanceId
14       AND bn.notification_type = 'OUT_OF_SERVICE'
15       AND bd.is_approved = 1
16 ORDER BY bn.notification_time, br.start_time;

```

§4.6 Requirement Traceability

Query	Required report	Principal relations
01	Approved hours for each space	Requests, decisions, spaces, space types
02	Approved count by weekday and hour	Decisions and requests
03	Capacity/facility/time room finder	Spaces, facilities, requests, decisions, maintenance
04	Bookings affected by escalation	Notifications, maintenance, decisions, requests, users

Table 9: Traceability of the four required analytical queries.

§5 Indexing and Query Tuning

§5.1 Scope and Methodology

The tuning experiment evaluates four workloads: the booking-conflict check, the Query 03 room finder, Query 01, and Query 02. The same generated database and query semantics were used before and after indexing.

The database ran Microsoft SQL Server 2022 at compatibility level 160. Each phase removed or created only the experimental indexes, refreshed statistics with `FULLSCAN`, warmed the workload, recompiled benchmark procedures, and measured each target five times. The complete baseline and tuned sequence was repeated three times. Logical reads, CPU time, elapsed time, access operators, index usage, and storage were recorded.

§5.2 Selected Indexes

Index	Workload rationale
IX_G08_BR_Space_Start	Keys <code>BOOKING.REQUEST</code> by <code>(space_code,start_time)</code> and includes end time/status for conflict and room-finder overlap checks.
IX_G08_BR_Start	Keys requests by <code>start_time</code> and covers space, end time, and status for Query 01 semester access and a narrower Query 02 scan.
IX_G08_BD_Approved_Booking	Filtered approved-only <code>booking_id</code> lookup for conflict checks and joins.
IX_G08_BD_Approved_DecisionTime	Filtered approved-only decision-time range with included booking identifier for Query 02.
IX_G08_Facility_Space_Name	Supports direct facility lookup by space and required facility name.
IX_G08_Maint_Space_Impact_Status_Start	Narrows maintenance by space, impact, active status, and interval start while covering the interval end.

Table 10: Final workload-specific index set.

§5.3 Before-and-After Results

Target	Logical reads		Reduction	Elapsed ms		Reduction
	Before	After		Before	After	
Conflict check	55.00	14.00	74.55%	0.31	0.08	73.12%
Room finder	35,152.20	953.40	97.29%	787.54	15.42	98.04%
Query 01	2,760.00	641.00	76.78%	42.31	28.73	32.10%
Query 02	2,677.00	829.00	69.03%	41.38	31.03	25.00%

Table 11: Average performance before and after indexing.




Figure 6: Elapsed-time comparison for the four tuned workloads.

§5.4 Execution-Plan Interpretation

Target	Before	After
Conflict check	Clustered request scan and decision lookup	Selective request and approved-decision seeks
Room finder	Repeated clustered scans of requests, facilities, and maintenance	Seeks on all repeated large/correlated accesses; the inexpensive 62-row space scan remains
Query 01	Clustered scans of both large booking relations	Semester range seek on requests and narrow approved-only decision scan
Query 02	Clustered scans of decisions and requests	Approved decision-time range seek and narrower covered request scan

Table 12: Important execution-plan changes.

The strongest change is the room finder: average elapsed time falls from 787.54 ms to 15.42 ms, while logical reads fall by 97.29%. Query 02 retains a residual `DATEPART(HOUR, decision_time)` expression, so its indexing gain is limited by grouping and computation after the semester range has been selected.

§5.5 Trade-offs and Recommendation

All six indexes were selected by SQL Server during measured runs. Together they use approximately 19.07 MB. The two request indexes account for most of the space and increase booking-write maintenance, but their key orders serve different access patterns: same-space interval checks versus semester ranges.

SQL Server also suggested a standalone request-status index. It was rejected because the condition `status <> 'cancelled'` qualifies roughly 93,029 of 100,030 requests and

is therefore not selective. Treating missing-index output as evidence rather than an automatic instruction avoids storage and write overhead with little expected benefit. The final recommendation is to retain the six measured workload-specific indexes, monitor their write cost and fragmentation, and reconsider them only if booking volume or report patterns change materially.

§6 Normalization Validation

§6.1 Functional Dependencies and Candidate Keys

The following dependencies are derived from the implemented identifiers and uniqueness constraints. The phrase *all non-key attributes* means every remaining descriptive attribute of that relation.

Table 13: Candidate keys and principal functional dependencies.

Relation	Candidate key(s)	Non-trivial functional dependencies
ROLE	role_id; role_name	role_id $\rightarrow$ role_name; role_name $\rightarrow$ role_id
USERS	user_id; email	user_id $\rightarrow$ all non-key attributes; email $\rightarrow$ all other attributes
SPACE_TYPE	space_type_id; space_type_name	space_type_id $\rightarrow$ space_type_name; space_type_name $\rightarrow$ space_type_id
SPACES	space_code; (building, room_number)	Each candidate key $\rightarrow$ all non-key attributes
AUTO_USAGE_POLICY	(space_type_id, role_id)	No non-key attributes; only trivial dependencies
FACILITY	facility_id	facility_id $\rightarrow$ space_code, facility_name, description
BOOKING_REQUEST	booking_id	booking_id $\rightarrow$ all non-key attributes
BOOKING_DECISION	decision_id; booking_id	decision_id $\rightarrow$ all non-key attributes; booking_id $\rightarrow$ all other attributes
USAGE_SESSION	session_id; decision_id	session_id $\rightarrow$ all non-key attributes; decision_id $\rightarrow$ all other attributes
MAINTENANCE_RECORD	maintenance_id	maintenance_id $\rightarrow$ all non-key attributes
BOOKING_NOTIFICATION	(booking_id, maintenance_id, notification_type)	Composite key $\rightarrow$ notification_time

Foreign keys do not create reverse functional dependencies. For example, a role can belong to many users and a space can have many requests. Conditional rules such as *automatic decisions have no deciding staff member* are business constraints, not additional functional dependencies.

§6.2 Normal-Form Proof

First Normal Form. All implemented columns contain atomic scalar values. Repeating sets such as required facilities and automatic policies are represented by separate rows and relations.

Second Normal Form. Relations with single-column keys have no possible partial dependency. `AUTO_USAGE_POLICY` has no non-key attributes. In `BOOKING_NOTIFICATION`, `notification_time` describes one complete booking–maintenance–notification-type event and depends on the full composite key.

Third Normal Form. Every non-trivial dependency listed in Table 13 has a candidate key as its determinant. No non-key attribute determines another non-key attribute within the same relation. Lookup descriptions such as role name and space-type name are stored only in their lookup relations, eliminating the Phase 1 transitive dependencies.

Consequently, every relation satisfies 3NF. Under the documented dependency set, every non-trivial determinant is also a superkey, so the relations satisfy the stronger Boyce–Codd Normal Form as well.

§7 LLM and Agent Improvement Process

§7.1 Model and Framework

The project used **DeepSeek V4 Flash** through the **OpenCode** framework. The LLM was used as a research and drafting assistant within a human-supervised process; final deliverables were reviewed and finalized by the group.

The framework separates reusable knowledge from experiment artifacts:

- the Global Agent stores project-wide terminology, invariants, consistency rules, and verification behavior;
- section-specific Skills store task methods and checklists;
- experiment results provide material for evaluation rather than becoming submissions directly; and
- human evaluation determines which reusable improvements are accepted.

§7.2 Improvement Cycle

1. Read the specification, current Agent knowledge, and relevant Skill.
2. Generate an experiment result for one deliverable.
3. Evaluate requirement coverage, naming consistency, SQL correctness, and unsupported assumptions.
4. Record the weakness, root cause, and reusable correction.
5. Refine the relevant knowledge and regenerate.
6. Human-review the final repository artifact.

For Phase 2, the review emphasis expanded in three directions. Concurrency claims required both an unsafe reproduction and a protected two-session result. Performance claims required repeatable measurements and execution-plan evidence rather than relying only on optimizer suggestions. Requirement traceability was extended from the ERD

through migration, maintenance notification, approval locking, and the four required analytical queries.

Evaluation concern	Reusable improvement
Concurrent correctness	State the invariant, reproduce the race, define one lock order, and verify the final database state.
Migration safety	Populate replacement columns and tables before removing Phase 1 structures; verify preserved counts.
Large-data validity	Make generation deterministic, namespace generated identifiers, and assert coverage/invariants inside the script.
Query performance	Compare identical semantics before and after indexing using repeated reads, timings, and plan operators.
Index discipline	Evaluate selectivity, storage, and write cost before accepting a missing-index recommendation.

Table 14: Representative Phase 2 agent-improvement lessons.

§8 Conclusion

Phase 2 successfully extends the Campus Space Management System for the operating conditions observed during the pilot. The schema now distinguishes advisory maintenance from out-of-service work, preserves maintenance-related booking notifications, normalizes roles and space types, unifies automatic and staff decisions, and migrates Phase 1 records without discarding their history.

The concurrency experiment demonstrates the central correctness risk and its prevention. The unsafe schedule produces two overlapping approvals, whereas the protected `BOOKING_REQUEST`–then–`SPACES` update-lock protocol leaves exactly one approved request and no overlapping approved pair.

The deterministic generator supplies three academic years and 100,000 new booking records. The four required reports cover utilization hours, weekday/hour distribution, suitable-space discovery, and maintenance-escalation impact. Six measured indexes improve every tuned workload; most notably, the room finder reduces logical reads by 97.29% and elapsed time by 98.04%.

Finally, the functional-dependency analysis shows that all eleven relations satisfy at least Third Normal Form. The resulting design provides a consistent basis for concurrent operations, operational reporting, and future growth while retaining clear trade-offs between read performance and index-maintenance cost.